

# NeuroPulse AI: System Audit & Analysis Report

Generated on: 2026-04-18 21:36:42

## 1. Executive Summary

---

NeuroPulse is a production-ready, highly-concurrent AI platform designed for deep emotional sentiment analysis and real-time misinformation detection. This document serves as a comprehensive forensic audit of the platform's architectural integrity, machine learning integration, and UI/UX compliance as of the final build.

## 2. Architectural Audit

---

The platform successfully implements a dual-branch, monolithic WSGI deployment bridging front-end react state management with complex server-side model routing.

- Backend (Flask): Acts as the robust orchestration layer running concurrent Python ThreadPoolExecutors. Highly effective request handling via '/analyze', '/fetch\_tweet', '/history', and '/stats'.
- Frontend (Vite + React): Responsive Single Page Application utilizing 'recharts' for telemetry mapping. Fully decoupled interface utilizing Glassmorphism design philosophies.
- Database (MongoDB Atlas): Centralized persistence layer maintaining execution traces and aggregate statistics via Document stores.

## 3. Model Engine Integrity

---

NeuroPulse boasts an unprecedented 11-Model operational matrix divided into two principal logic branches:

### Branch A: The Sentiment Ensemble (10 Models)

- Classic ML (5 Models): SVM, Logistic Regression, Random Forest, XGBoost, Naive Bayes. Successfully trained via 'train\_sentiment.py' wrapping a 10k max-feature TF-IDF Vectorizer.
- Deep Learning (2 Models): PyTorch BiLSTM and 1D-CNN. Weights stored properly locally.
- Transformers (3 Models): DistilBERT, BERT, RoBERTa. Directly fetched from HuggingFace cache dynamically.

### Branch B: The Misinformation Classifier

- Utilizes a completely independent XGBoost classifier mapping exclusively for fake-news markers.
- Trained autonomously via 'train\_misinformation.py' executing against a pristine uniformly-distributed 30,000-row synthetically harvested dataset ('Misinformation\_Data.csv').
- Fallback OOD catchers active for specific non-vectorized high-risk inputs (e.g. 5G Microchips).

## 4. Gemini Search Grounding Verification

---

A revolutionary addition to the platform pipeline. Rather than executing an LLM arbitrarily, Gemini 2.5 Flash

# NeuroPulse AI: System Audit & Analysis Report

Generated on: 2026-04-18 21:36:42

operates strictly defensively:

- Triggered strictly via the XGBoost 'High' risk label.
- Prompt structure rigorously enforced to extract actionable intelligence ('FINDING: ...').
- Employs Google Search Grounding to anchor outputs to live factual databases, successfully escaping hallucination risks.
- Validates evidence with verifiable URL extractions mapped to the frontend component 'FactCheckEvidence.jsx'.

## 5. Security & Stability Enhancements

---

- React UI state robustness: Enforced widespread Optional Chaining to prevent undefined references (e.g., 'item?.sentiment?.toUpperCase()') from crashing the History interface.
- Graceful ML Degradation: Local OS '.pkl' and '.pt' loading uses try-except blocks preventing full systemic failure if an individual model corrupts.
- Conflict Anomaly Module: Cross-validation between the Sentiment Ensemble and Misinformation matrix intelligently spots AI divergence (Positive Sentiment + High Misinformation), notifying human reviewers.

## 6. Conclusion & Deployment Checklist

---

NeuroPulse has successfully cleared Phase 4 UI/UX Verification and Phase 5 Integrations. The platform stands functionally complete. Recommended production deployment steps:

- 1. Migrate Flask from built-in dev-server to Gunicorn/Waitress.
- 2. Bundle React assets via 'npm run build'.
- 3. Configure system environment variables (GEMINI\_API\_KEY, MONGO\_URI) natively into the cloud deployment instance.
- 4. Transition into containerized runtime (Docker) depending on hardware acceleration needs for Torch models.